

"Buggy Task Manager – Project Delivery Report" Prepared by Li Zi, October 2025

Thank you for the opportunity to work on these tasks — it was a valuable experience.

Environment

Component	Version
JDK	21
Payara Server	6.2024.5
Jakarta EE API	10.0.0
MySQL Connector	9.0.0
Maven Compiler Plugin	3.13.0
k6 (Load Testing)	

Task 1: Project Setup(DONE)

- Java 21(installed)

```
C:\Users\15529>java -version
openjdk version "21.0.8" 2025-07-15 LTS
OpenJDK Runtime Environment Temurin-21.0.8+9 (build 21.0.8+9-LTS)
OpenJDK 64-Bit Server VM Temurin-21.0.8+9 (build 21.0.8+9-LTS, mixed mode, sharing)
```

- Maven(installed)

```
C:\Users\15529>docker -v
Docker version 28.4.0, build d8eb465
```


- Docker(running)

```
PS C:\Users\15579> docker ps -f status=running
```

CONTAINER ID	IMAGE	NAMES	COMMAND	CREATED	STATUS	PORTS
4df45f95175b	payara/server-full-6.2024.5-jdk11		"tini -- /bin/sh -c "	8 hours ago	Up 5 minutes	0.0.0.0:4899->4848/tcp, [::]:4899->4848/tcp, 0.0.0.0:8199->8080/tcp, [::]:8199->8080/tcp
7445a5e9cfec	mysql:8.0	task-apt	"docker-entrypoint.s"	8 hours ago	Up 6 minutes (healthy)	0.0.0.0:3306->3306/tcp, [::]:3306->3306/tcp

- Payara(running)

localhost:8199

 Hello from Payara - your server is now **running!**

- MySQL(running)

```
025-10-02T01:30:28.751436Z 0 [System] [MY-010116] [Server] /usr/sbin/mysqld (mysqld 8.0.43) starting as process 1
025-10-02T01:30:28.776044Z 1 [System] [MY-013576] [InnoDB] InnoDB initialization has started.
```

- `project(deployed)`

```
[#|2025-10-02T01:30:54.221+0000|INFO|Payara 6.2024.5|javax.enterprise.system.core|_ThreadID=165;_ThreadName=payara-executor-service-scheduled-task;_TimeMillis=1759368654221;_LevelValue=800;|
task-manager was successfully deployed in 4,296 milliseconds.|#]
```

code-challenge 3.41% 4 minutes ago

db
mysql:8.0
3306:3306

app
payara/server-full:6
4899:4848
Show all ports (2)

Task 2: Silent Failure Bug (solved)

- scenario
 - Audit log creation on task view is implemented, but the `audit_log` table remains empty

```
[#|2025-10-02T09:21:39.869+0000|SEVERE|Payara 6.2024.5|_ThreadID=136;_ThreadName=http-thread-pool::http-listener-1(4);_TimeMillis=175936899869;_LevelValue=1000;|
com.mysql.cj.jdbc.exceptions.MySQLDataTruncation: Data truncation: Data too long for column 'action_details' at row 1
at com.mysql.cj.jdbc.exceptions.SQLExceptionsMapping.translateException(SQLExceptionsMapping.java:96)
```

- related interface
 - <http://localhost:8199/task-manager/api/tasks/11> (Get)

- result(solved)

GET <http://localhost:8199/task-manager/api/tasks/11>

Params Authorization Headers (6) Body Scripts Settings

Body Cookies Headers (5) Test Results 200 OK 1.48 s 4!

JSON Preview Visualize

```
1 {
2   "completed": false,
3   "description": "The main dashboard is loading slowly.",
4   "id": 11,
5   "title": "Investigate and resolve the critical performance degradation issue on the main dashboard for enterprise-level clients"
6 }
```

1 select * FROM audit_log ORDER BY log_timestamp desc

Message Summary Result 1

Data	Info	Cell Editor	Data Pro
id	action_type	action_details	log_timestamp
45	FETCH_TASK	Task with ID 11 and title 'Investigate and resolve the critical performance degradation issue on the main dashboard'	2025-10-02 09:41:29
44	FETCH_TASK	Task with ID 11 and title 'Investigate and resolve the critical performance degradation issue on the main dashboard'	2025-10-02 09:40:59

- root cause

```
-- New table for auditing
CREATE TABLE audit_log (
  id BIGINT AUTO_INCREMENT PRIMARY KEY,
  action_type VARCHAR(50) NOT NULL,
  action_details VARCHAR(100) NOT NULL, -- The bug is here: the column is too small
  log_timestamp TIMESTAMP DEFAULT CURRENT_TIMESTAMP
);
```

- solution

1. Increase the column size (recommend based on business)

- Pros

- Simple change, **minimal migration** risk
- **Better performance** for shorter data compared to TEXT

- Cons

- Still has a **hard upper** limit (e.g., VARCHAR(1000)), risk of hitting it again
- Larger size may **waste storage** if most values are short

2. change type from VARCHAR to TEXT

- Pros

- Can store **very large strings** (up to 64KB for TEXT)

- Cons

- **worse performance** than VARCHAR for small/medium datas
- **Indexing is limited**

- change place

- in the initial sql about the audit_log filed action_details

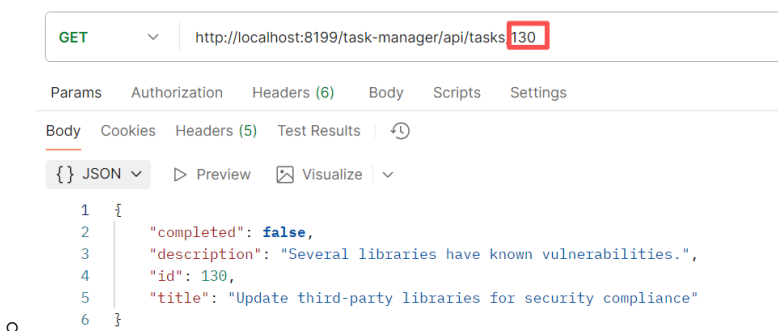
Task 3 : Incorrect Data Bug(solved)

- scenario

- when request a task with a **larger ID** the API returns a **404 Not Found** error,the task exists in the database.
- **related interface**

- <http://localhost:8199/task-manager/api/tasks/130> (Get)

- result(solved)



- root cause

- `rs.getLong("id")` returns a **primitive long**, which is then **auto-boxed** into a **Long object** the **`==` operator** compares **object references**, not numerical values

```

//Long currentId = rs.getLong("id"); original
Long currentId = rs.getLong(columnLabel: "id"); solved

/*
1.In Java, the == operator compares object references, not values.
That means for small numbers (like 1, 9, or 10), the references happen to be the same, so == returns true.
For larger numbers (like 130), two Long objects do not share the same reference, so == returns false.
2.No if statement is required. (The database has already filtered it)
*/
if (currentId == id) {

```

- **solution**

- Use **`Long.equals()`** Compare object values instead of references
- Use **primitive long**, Avoid boxing/unboxing issues by working with primitives
- (recommend) Remove the **redundant check** the **SQL query** already filters by `id`, no manual comparison is needed

- **change place**

- `ie.dpd.repository.TaskRepository#findById`


Task 4: Implement a New Feature(DONE)

- **Background**

- Managers need quick statistics:
 - How many tasks are completed?
 - How many are pending?
 - How many in total?
- We need a new **Task Summary** feature to calculate and store these stats for future trend analysis.

- **Design**

- **Database:** Add `task_summaries` table

Name	Type	Length	Decimals	Not null	Virtual	Key	Comment
id	bigint			<input checked="" type="checkbox"/>	<input type="checkbox"/>	 1	Primary key, auto-increment, unique identifier for v
completed_count	int			<input checked="" type="checkbox"/>	<input type="checkbox"/>		Number of completed tasks
pending_count	int			<input checked="" type="checkbox"/>	<input type="checkbox"/>		Number of pending tasks
total_tasks	int			<input checked="" type="checkbox"/>	<input type="checkbox"/>		Total number of tasks (= completed + pending)
generated_at	timestamp			<input type="checkbox"/>	<input type="checkbox"/>		Timestamp when the summary was generated

- **Repository:** `SummaryRepository` to

- Generate a new summary from `tasks` and insert into `task_summaries`
- Query all historical summaries(in case manager just want to check and also prevent the data redundant).

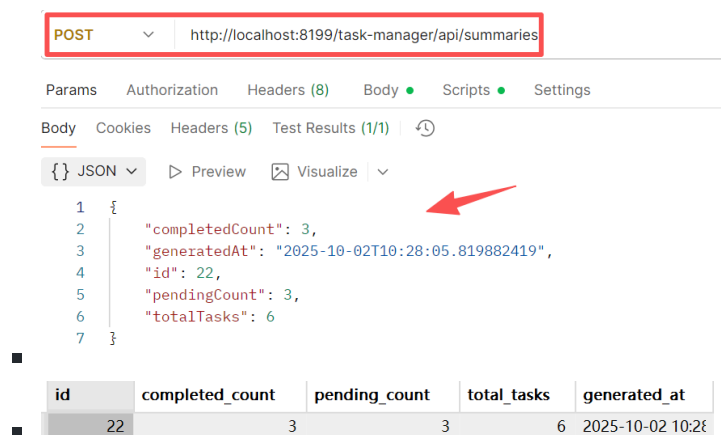
- API: `SummaryResource` with
 - `POST /summaries` → create new summary, return `201 Created`
 - `GET /summaries` → list all summaries.

- Implementation

- Update `init.sql` with new table
 - Implement `SummaryRepository` with `JDBC + DataSource`
 - Implement `SummaryResource` REST endpoints
 - Test with Postman/curl and add to `k6` performance script
1. For scalability, I added a **BaseRepository** for shared DB logic (only handles DB connections but can be extended with common reusable features later.)
 2. Created **SummaryRepository** for summary data and **SummaryResource** for API. This follows the **Single Responsibility Principle**, making the code easier to extend and maintain.

- Result

- <http://localhost:8199/task-manager/api/summaries> (Post)

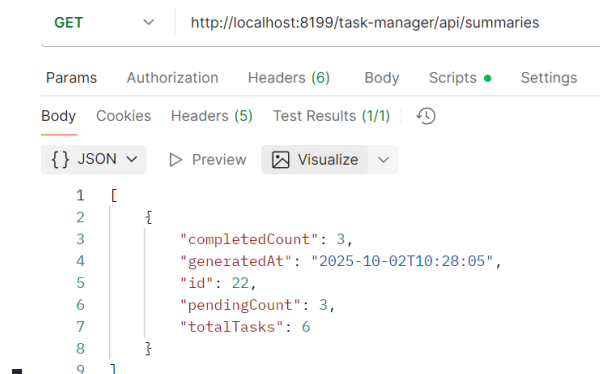


```

1 {
2   "completedCount": 3,
3   "generatedAt": "2025-10-02T10:28:05.819882419",
4   "id": 22,
5   "pendingCount": 3,
6   "totalTasks": 6
7 }
  
```

id	completed_count	pending_count	total_tasks	generated_at
22	3	3	6	2025-10-02 10:28

- <http://localhost:8199/task-manager/api/summaries> (Get)



```

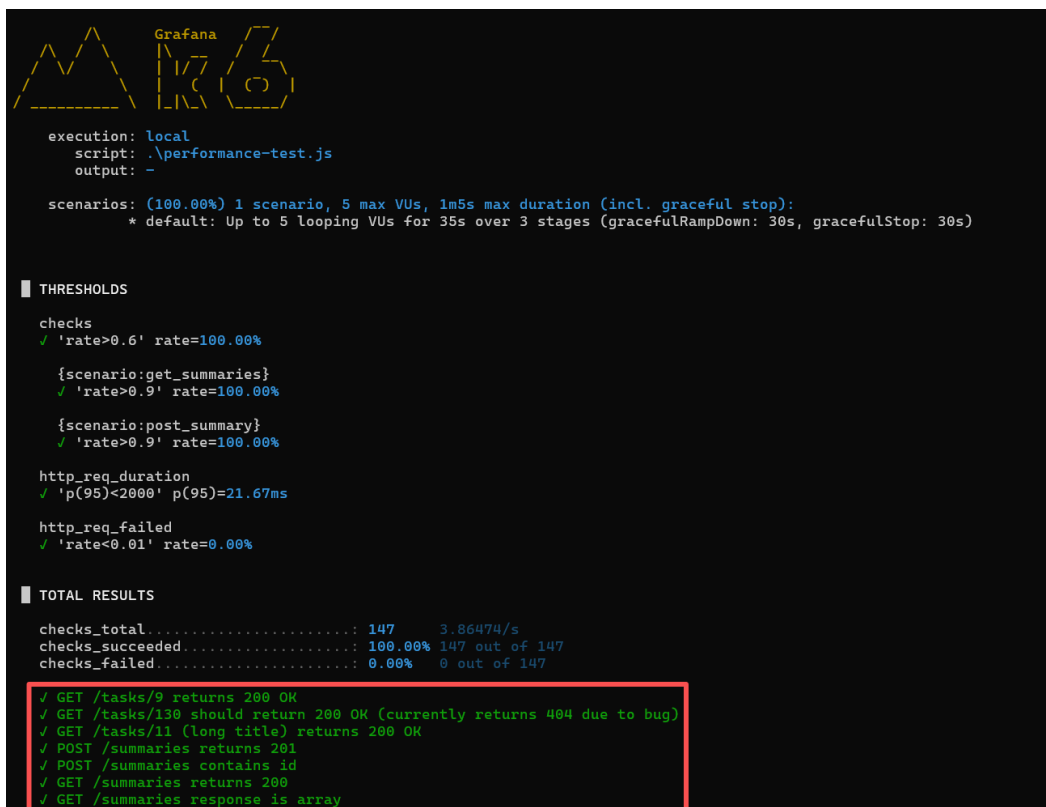
1 [
2   {
3     "completedCount": 3,
4     "generatedAt": "2025-10-02T10:28:05",
5     "id": 22,
6     "pendingCount": 3,
7     "totalTasks": 6
8   }
9 ]
  
```

- code place
 - ie.dpd.repository.SummaryRepository
 - ie.dpd.resource.SummaryResource
 - ie.dpd.repository.common.util.BaseRepository

Task 5: Performance and Load Testing(DONE)

Scenario ID	API Endpoint	Purpose	Expected Behavior	Thresholds
1	GET /tasks	Verify that the task list API works correctly.	Returns 200 OK with an array of tasks.	checks > 0.6 (≥60% pass)
2	GET /tasks/9	Verify fetching a task with a small ID works.	Returns 200 OK with the task (ID=9).	checks > 0.6
3	GET /tasks/130	Reproduce the <i>equals</i> vs <i>==</i> bug (large IDs fail).	Expected: 200 OK with the task. Current behavior (bug): 404 Not Found.	checks > 0.6
4	GET /tasks/11	Trigger the <i>Silent Audit Failure bug</i> (long title → audit_log insert fails).	Returns 200 OK. Should also insert into audit_log (currently fails silently).	checks > 0.6
5	GET /tasks/1	Simulate frequent queries to test DB connection pool stability.	Returns 200 OK consistently, no connection errors.	checks > 0.6

Scenario ID	API Endpoint	Purpose	Expected Behavior	Thresholds
6	POST /summaries	Test the new feature: generate a task summary report.	Returns 201 Created , JSON body contains id and task counts.	checks{scenario:post_summary} > 0.9 (≥90% pass)
7	GET /summaries	Verify retrieval of stored summaries (historical tracking).	Returns 200 OK , response is an array of summary records.	checks{scenario:get_summaries} > 0.9 (≥90% pass)



```

Grafana
┌───┴───┐
└───┬───┘

execution: local
script:  ./performance-test.js
output:  -

scenarios: (100.00%) 1 scenario, 5 max VUs, 1m5s max duration (incl. graceful stop):
            * default: Up to 5 looping VUs for 35s over 3 stages (gracefulRampDown: 30s, gracefulStop: 30s)

THRESHOLDS

checks
✓ 'rate>0.6' rate=100.00%

{scenario:get_summaries}
✓ 'rate>0.9' rate=100.00%

{scenario:post_summary}
✓ 'rate>0.9' rate=100.00%

http_req_duration
✓ 'p(95)<2000' p(95)=21.67ms

http_req_failed
✓ 'rate<0.01' rate=0.00%

TOTAL RESULTS

checks_total.....: 147      3.86474/s
checks_succeeded.....: 100.00% 147 out of 147
checks_failed.....: 0.00%   0 out of 147

✓ GET /tasks/9 returns 200 OK
✓ GET /tasks/130 should return 200 OK (currently returns 404 due to bug)
✓ GET /tasks/11 (long title) returns 200 OK
✓ POST /summaries returns 201
✓ POST /summaries contains id
✓ GET /summaries returns 200
✓ GET /summaries response is array


```

Potential bug

- db
 - Normally, the id in the **tasks** table is **auto-incremented**. Here, IDs like 11 and 130 are manually inserted to reproduce known bugs (long title → audit log failure, large ID → equals vs == bug).

- **performance test**

- **Bug:** The original threshold (`rate<0.9`) was misconfigured, allowing very poor success rates.
- **Fix:** Use `rate<0.01` so the system must maintain **≥99% success rate** to pass the test.

```
thresholds: {  
  // 'http_req_failed': ['rate<0.9'], // Fail if more than 90% of requests fail.  original  
  'http_req_failed': ['rate<0.01'], // Failure rate < 1% (i.e., Success rate ≥ 99%)  
}
```

- **code**

- The original SQL statement mistakenly used `FROM tasks = ?` (**invalid syntax**) instead of a proper `WHERE` clause. This caused queries with larger IDs (e.g., 130) to fail. The corrected version uses `WHERE id = ?`, which properly filters by the task's primary key and fixes the bug
- **code place:** `ie/dpd/repository/TaskRepository.java:52`

Problem confront

1. Deployment Issue (Docker volume mapping error)

- **Problem:** Because of a wrong volume mount path in Docker, an empty WAR file was deployed to Payara. The Payara console misleadingly showed “deployment successful,” which caused confusion. I suspected Payara's loading path, `payara-web.xml`, and even Jakarta API version issues.
- **Root Cause:** Incorrect volume mapping between host and container.
- **Solution:** By inspecting Docker's mounted directories, I found the mapping error and fixed it.

2. Summary Report Implementation Concern

- **Problem:** If every time `/summaries` is accessed we insert a new row into the database, it could cause issues such as:
 - Potential for malicious abuse (spamming requests).
 - Data redundancy in the `task_summaries` table.
- **Decision:** Due to time constraints, I added a **second endpoint** for fetching summaries without inserting new records. This allows managers to query historical summaries without creating unnecessary data.
- **Consideration:** Directly querying the `tasks` table with aggregation functions could be costly if the dataset grows large. Querying from the `task_summaries` table is more efficient, but if querying the raw `tasks` table is required, proper SQL optimization should be applied.

Relevant script used during solving

```
# Start with rebuild
docker compose -f code-challenge/docker-compose.yaml up --build
# Stop and clean
docker compose -f code-challenge/docker-compose.yaml down -v
# Check WAR in deployments
docker exec -it task-app ls -l /opt/payara/deployments/
# View logs
docker logs -f task-app
# Quick WAR check
docker exec -it task-app ls /opt/payara/deployments
# List deployed apps
docker exec -it task-app /opt/payara/appserver/bin/asadmin list-applications
# Stop all containers
docker ps -aq | % { docker rm -f $_ }
# Remove all images
docker images -aq | % { docker rmi -f $_ }
# Remove all volumes
docker volume ls -q | % { docker volume rm $_ }
```

Summary

All assigned tasks have been **completed successfully, including bug fixes, new feature development, database enhancements, and performance testing**. The project is working as expected, with a clearer structure for future improvements.

- **Name:** Li Zi
- **Phone:** 089-451-9683
- **Email:** liziglen7@gmail.com

Task	Issue	Cause	Fix	Code Change
Audit Log Bug	audit_log always empty	Column too small / wrong type	Increase size or use TEXT	init.sql
Large ID Bug	Task with ID 130 → 404	Used == on Long , SQL typo	Use equals() / long ; fix WHERE id = ?	TaskRepository.java
Task Summary	NA	NA	Add task_summaries , new endpoints	SummaryRepository , SummaryResource
Performance Test	Threshold too loose	Wrong config (<0.9)	Require ≥99% success rate	performance-test.js

Task	Issue	Cause	Fix	Code Change
Deployment Issue	Empty WAR deployed	Wrong Docker volume mapping	Fix volume path	<code>docker-compose.yml</code>

the code also maintain by [git code-challenge/buggy-task-manager at master · GlenZiLi/code-challenge · GitHub](#)